

Introduction to Julia: Why are we doing this to you?

(Spring 2019)

Steven G. Johnson, MIT Applied Math

MIT classes 18.06, 18.303, 18.330, 18.08[56],
18.335, 18.337, ...

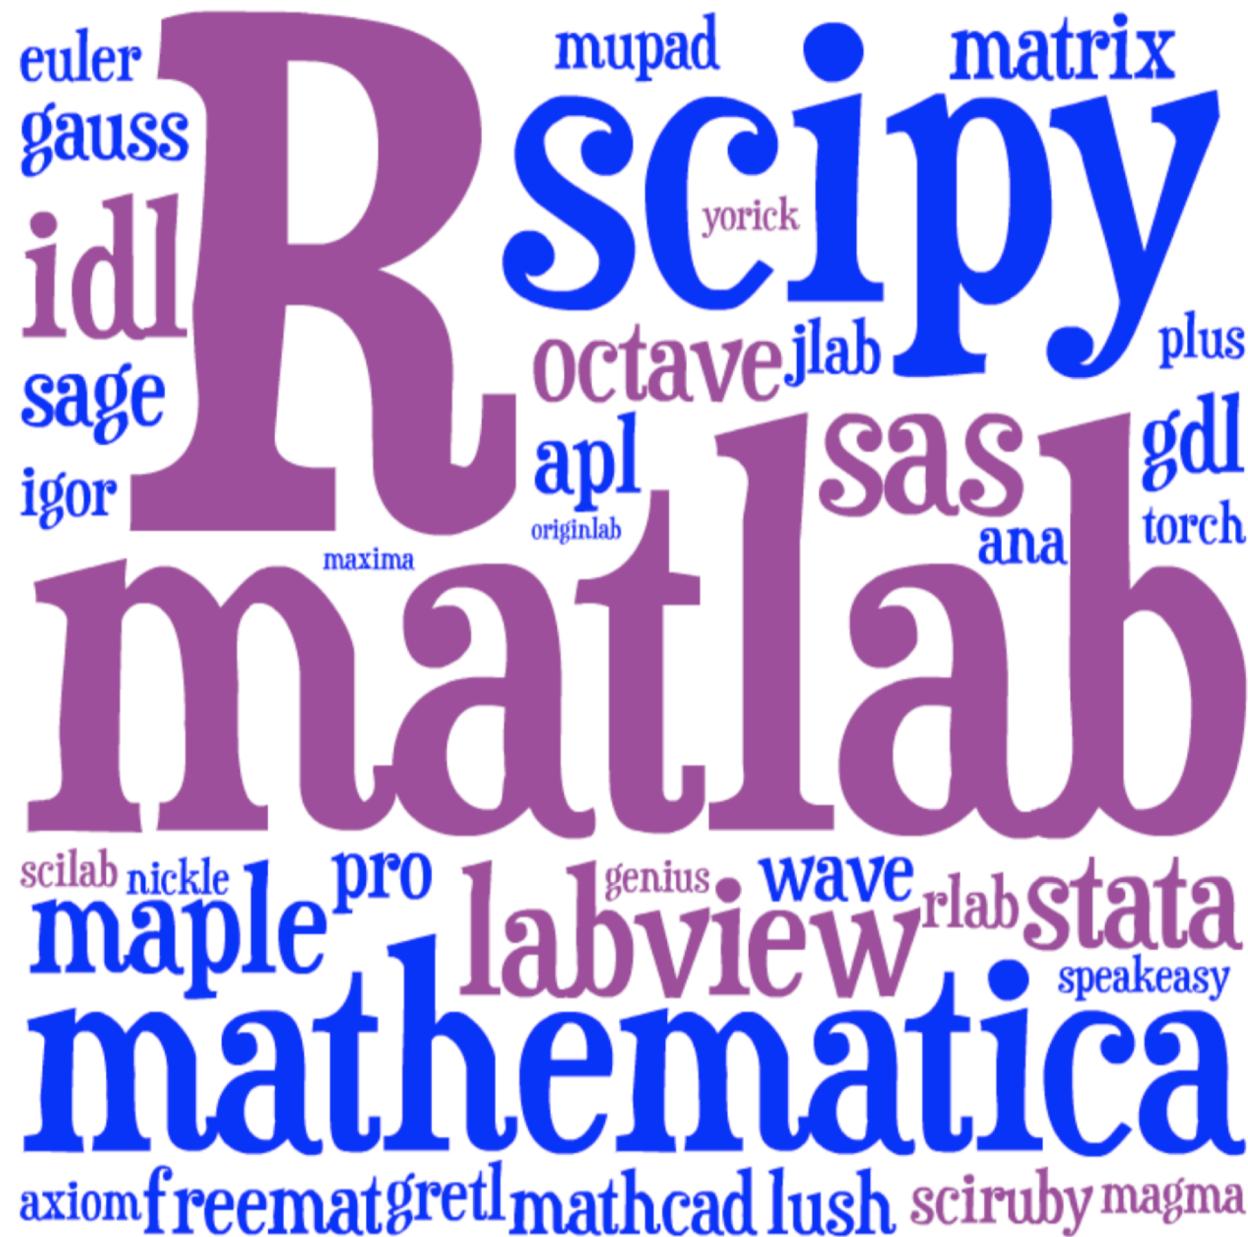
What language for teaching scientific computing?

For the most part, these are not hard-core programming courses, and we only need little “throw-away” scripts and toy numerical experiments.

Almost any high-level, interactive (dynamic) language with easy facilities for linear algebra ($Ax=b$, $Ax=\lambda x$), plotting, mathematical functions, and working with large arrays of data would be fine.

And there are lots of choices...

Lots of choices for interactive math...



[image: Viral Shah]

Just pick the most popular?
Matlab or Python or R?

*We feel guilty pushing a language on
you that we
are starting to abandon ourselves.*

Traditional HL computing languages
hit a performance wall in “real” work
... eventually force you to C, Cython, ...

A new programming language?

Viral Shah

Jeff Bezanson



Alan Edelman



[MIT]

Stefan Karpinski

[30+ developers with 100+ commits,
1000+ external packages, 6th JuliaCon in 2019]



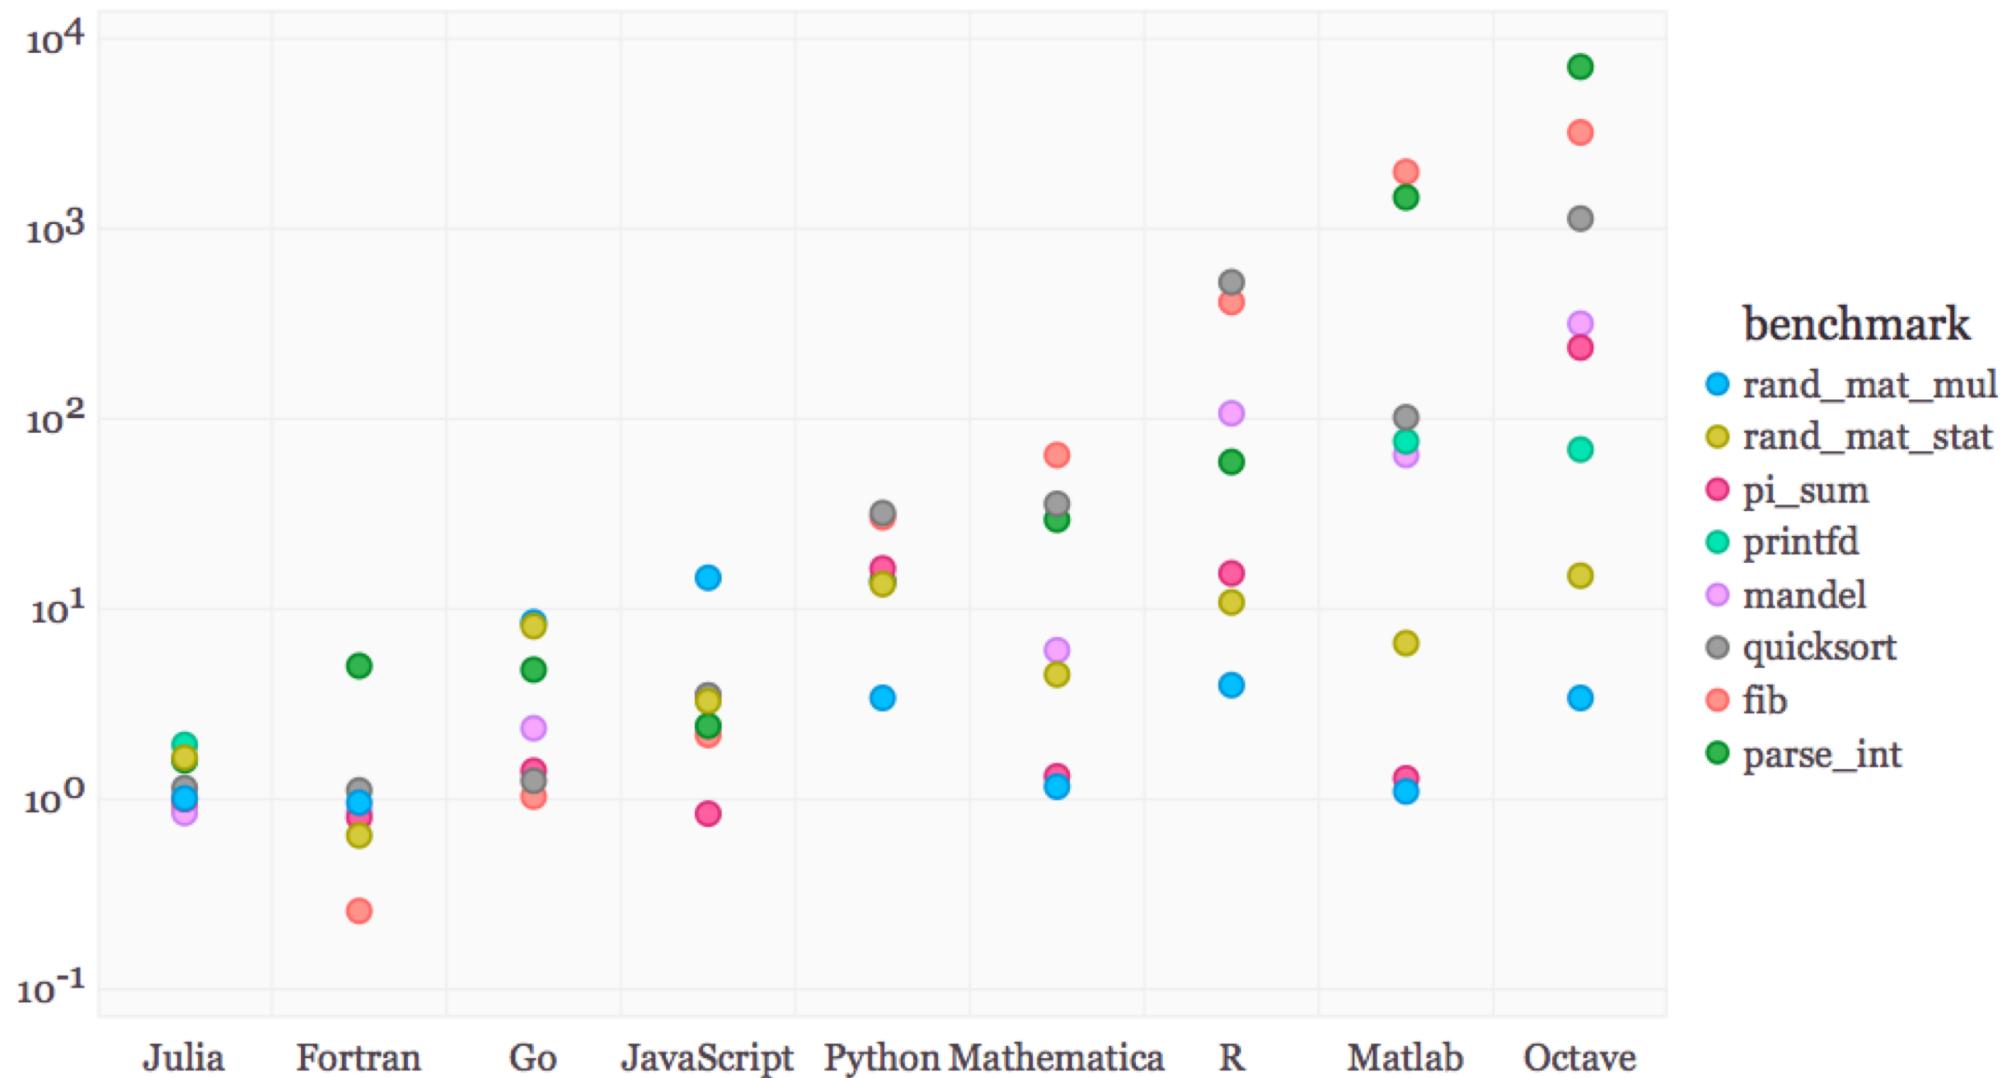
julialang.org

[begun 2009, “0.1” in 2013, ~40k commits,
1.0 release in Aug. 2018, 1.1 in Jan. 2019]

As **high-level and interactive** as Matlab or Python+IPython,
as **general-purpose** as Python,
as productive for **technical** work as Matlab or Python+SciPy,
but as **fast as C**.

Performance on synthetic benchmarks

[loops, recursion, etc., implemented in most straightforward style]



(normalized so that C speed = 1)

Special Functions in Julia

Special functions $s(x)$: classic case that cannot be vectorized well

... switch between various polynomials depending on x

Many of Julia's special functions come from the usual C/Fortran libraries,
but some are written in pure Julia code.

Pure Julia `erfinv(x)` [= $\text{erf}^{-1}(x)$]

3–4× faster than Matlab's and 2–3× faster than SciPy's (Fortran Cephes).

Pure Julia `polygamma(m, z)` [= $(m+1)^{\text{th}}$ derivative of the $\ln \Gamma$ function]

~ 2× faster than SciPy's (C/Fortran) for real z

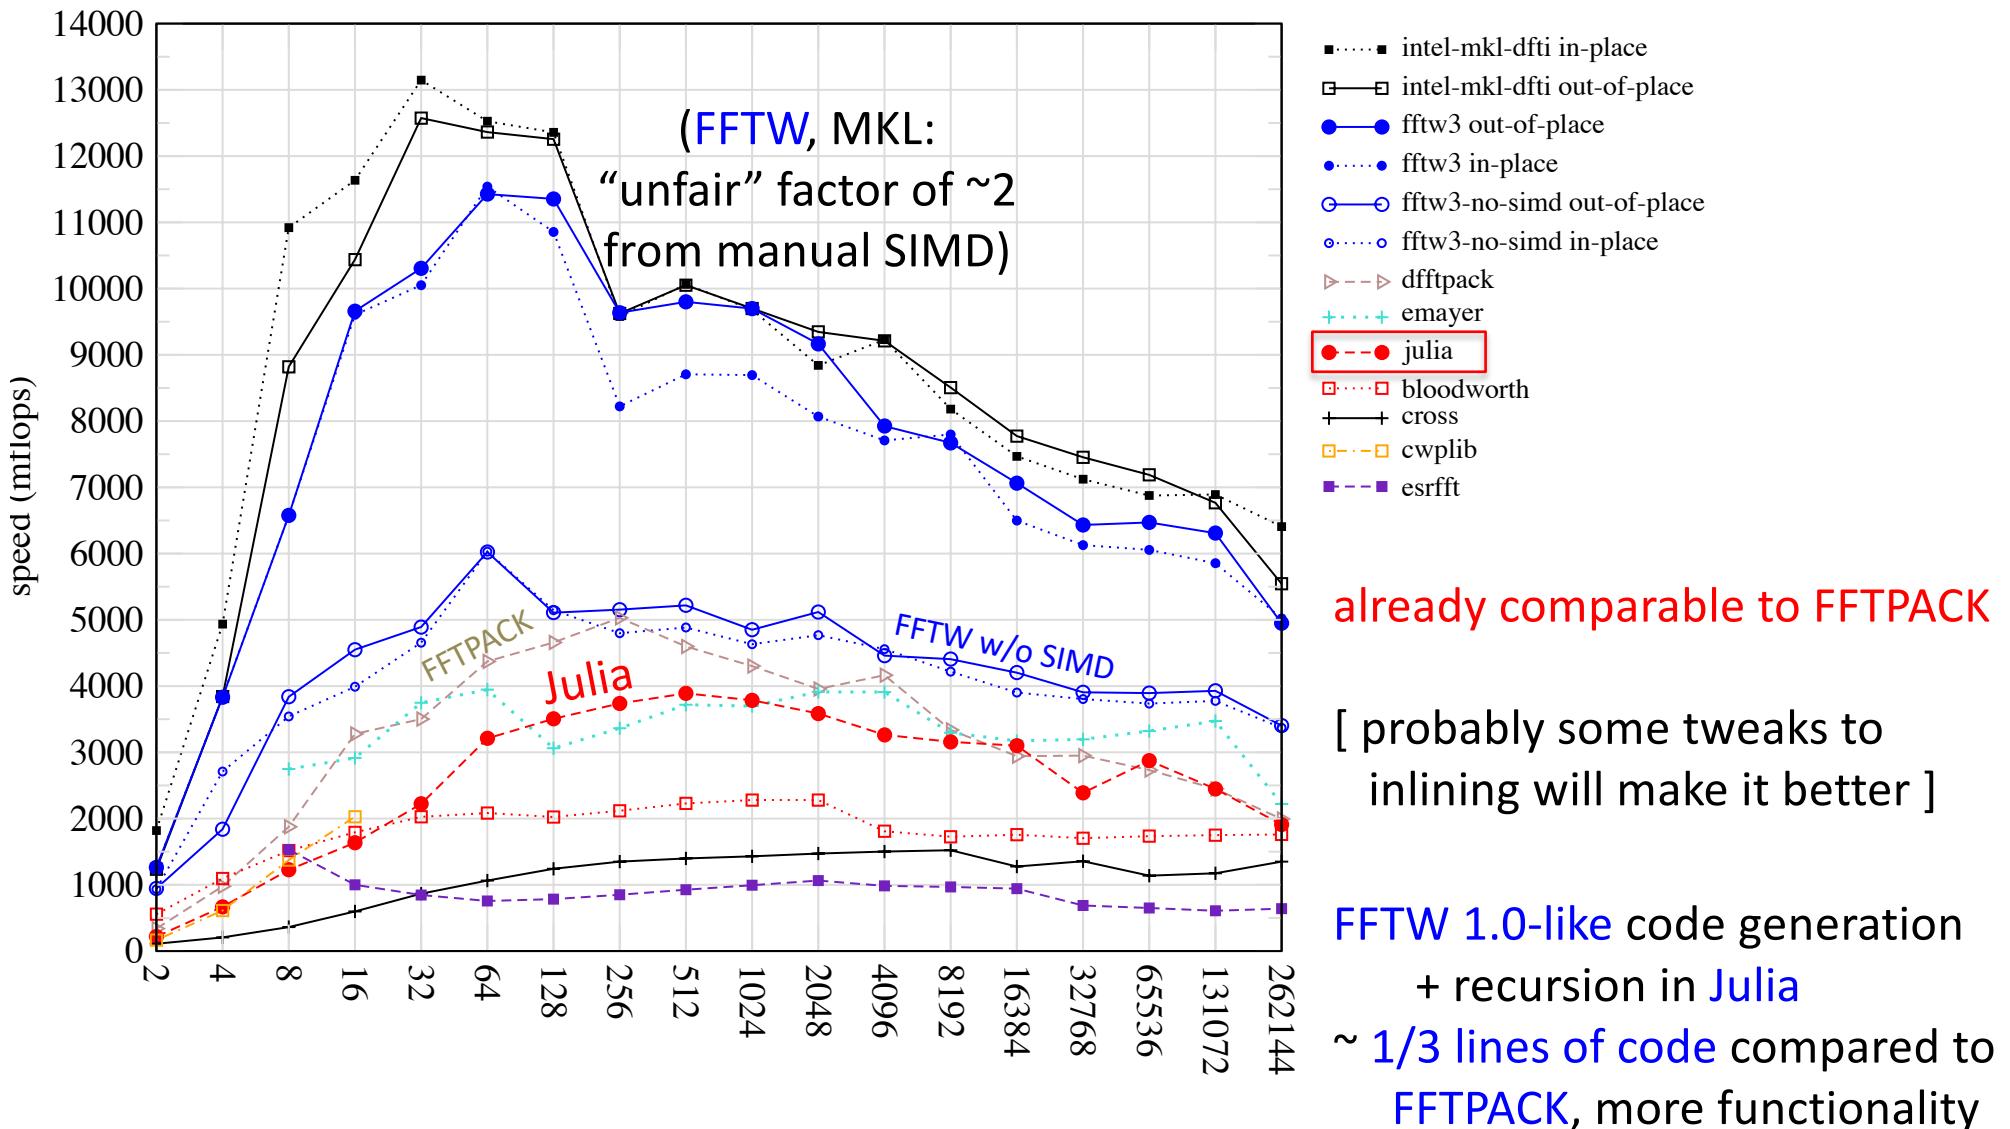
... and unlike SciPy's, same code supports complex argument z

Julia code can actually be faster than typical “optimized”
C/Fortran code, by using techniques
[metaprogramming/codegen generation] that are
hard in a low-level language.

Pure-Julia FFT performance

double-precision complex, 1d transforms

powers of two



Generating Vandermonde matrices

given $x = [\alpha_1, \alpha_2, \dots]$, generate:

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

NumPy (`numpy.vander`): [follow links]

[Python code](#) ...wraps [C code](#)
... wraps [generated C code](#)

type-generic at high-level, but
low level limited to small set of types.

Writing fast code “in” Python or Matlab = [mining the standard library](#)
for pre-written functions (implemented in C or Fortran).

If the problem doesn’t “vectorize” into built-in functions,
if you have to write your [own inner loops](#) ... [sucks](#) for you.

Generating Vandermonde matrices

given $x = [\alpha_1, \alpha_2, \dots]$, generate:

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

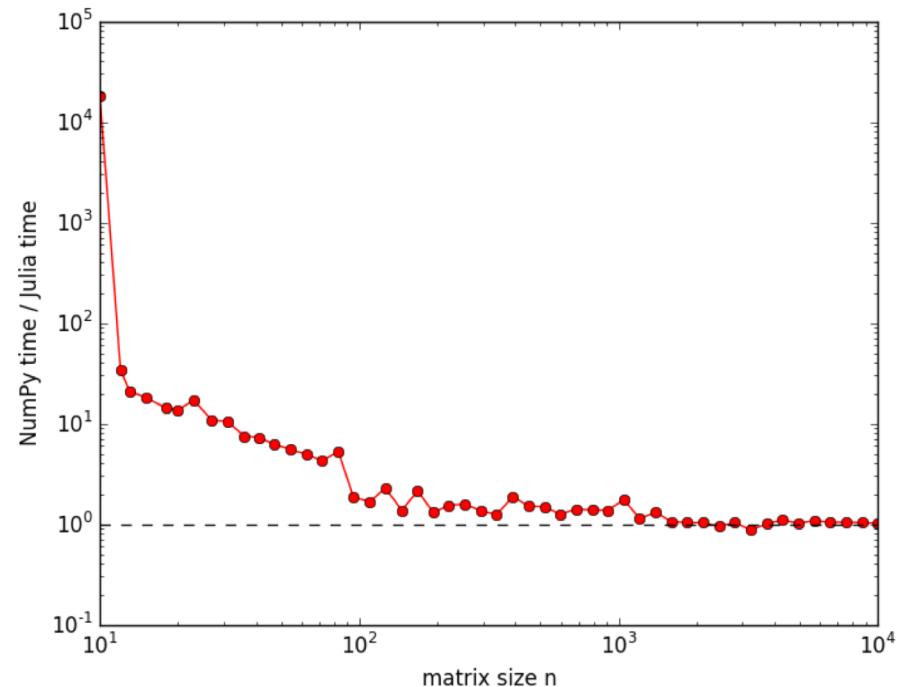
Julia (type-generic code):

```
function vander(x, n=length(x))
    m = length(x)
    V = Array(eltype(x), m, n)
    for j = 1:m
        V[j,1] = one(x[j])
    end
    for i = 2:n
        for j = 1:m
            V[j,i] = x[j] * V[j,i-1]
        end
    end
    return V
end
```

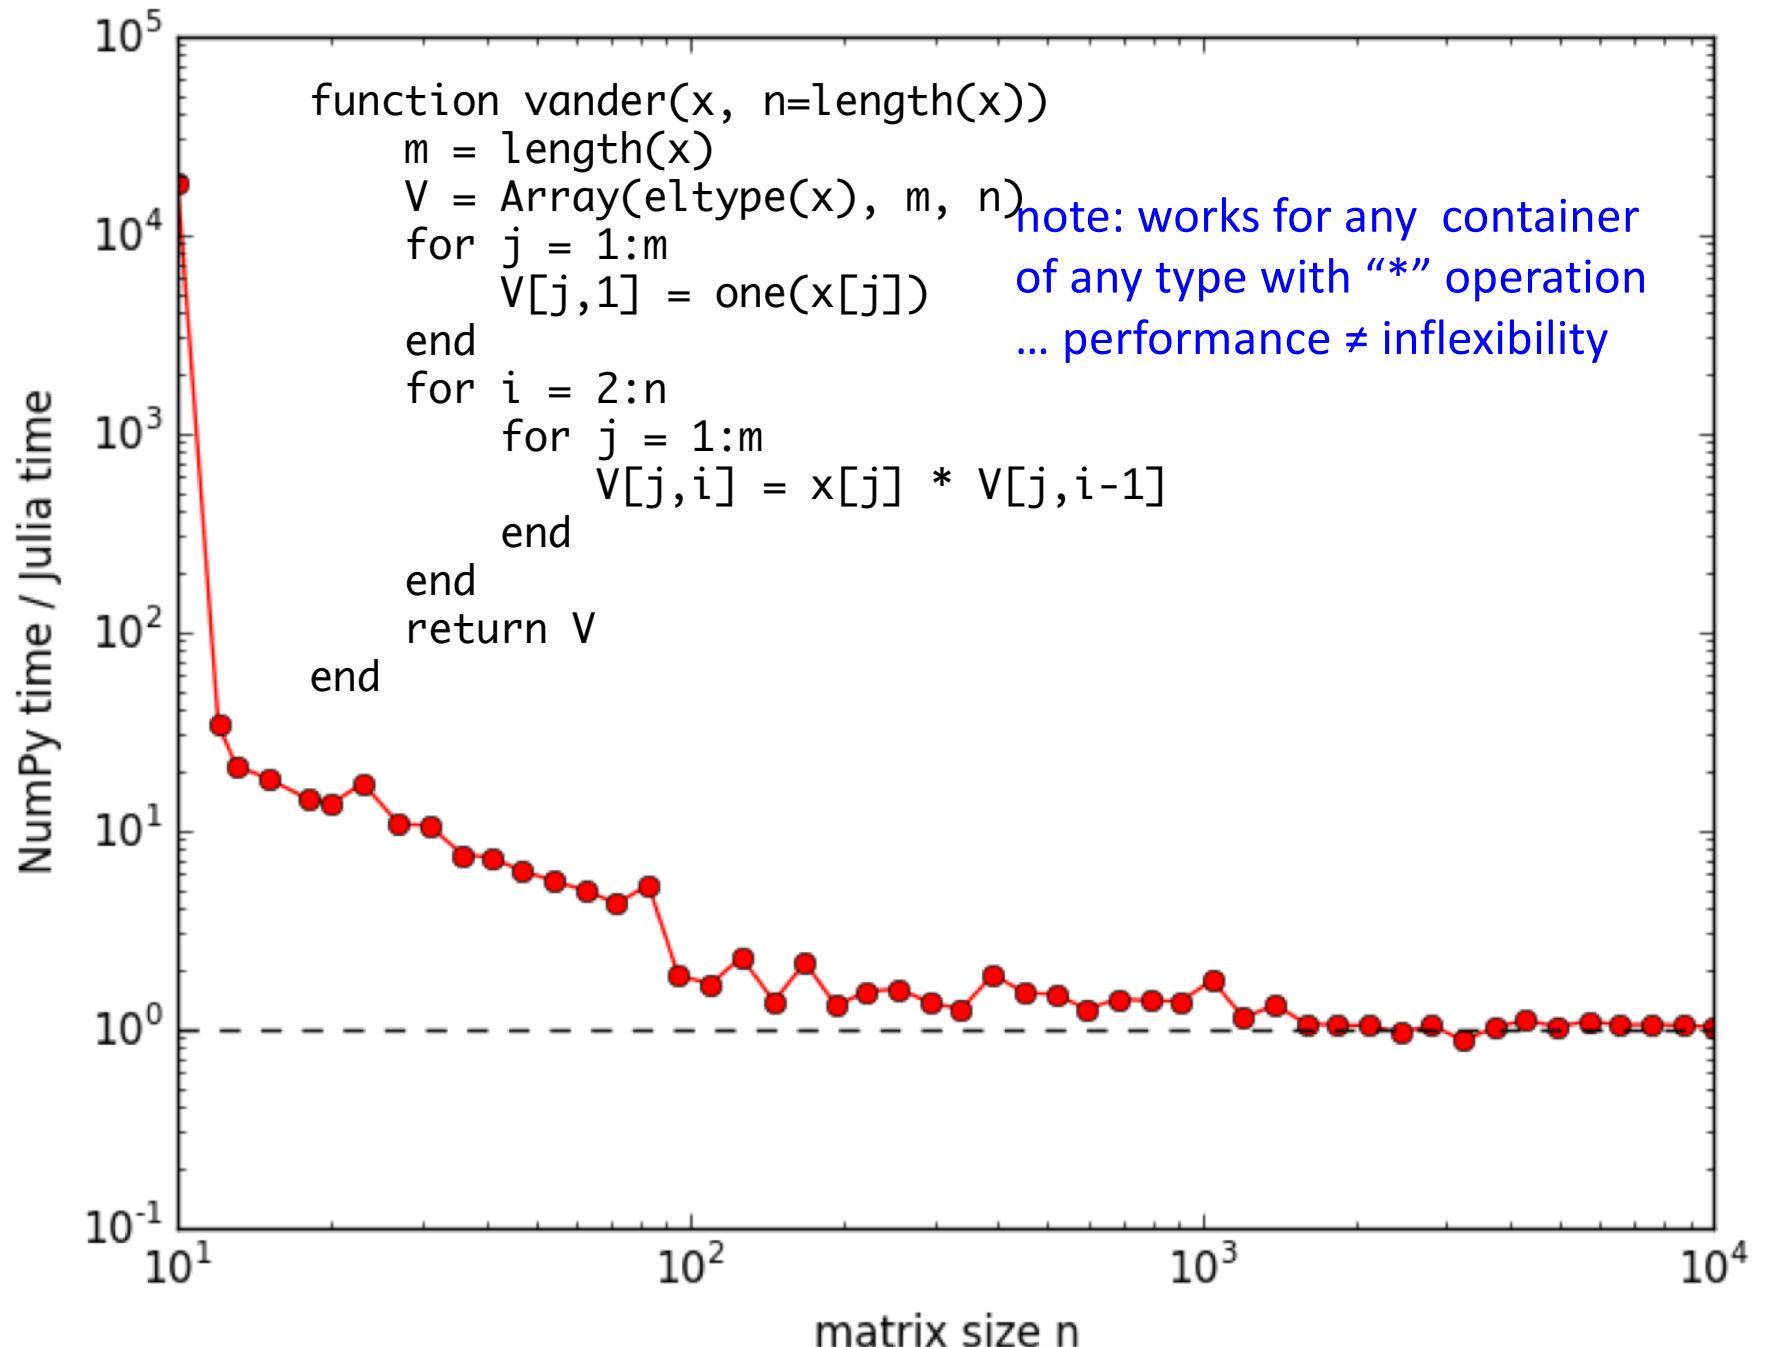
NumPy (`numpy.vander`): [follow links]

[Python code](#) ...wraps [C code](#)
... wraps [generated C code](#)

type-generic at high-level, but
low level limited to small set of types.



Generating Vandermonde matrices



But I don't “**need**” performance!

For lots of problems, especially “toy” problems in courses, Matlab/Python performance is **good enough**.

But if use those languages for all of your “easy” problems, then **you won't be prepared to switch when you hit a hard problem**. When you **need** performance, it is too late.

You **don't** want to learn a **new language** at the **same time** that you are solving **your first truly difficult** computational problem.

Just vectorize your code?

= rely on mature **external libraries**,
operating on **large blocks of data**,
for performance-critical code

Good advice! But...

- **Someone** has to write those libraries.
- Eventually that person will be **you**.
 - **some problems** are impossible or just very awkward to vectorize.

But everyone else is using Matlab/Python/R/...

Julia is still a young, niche language. That imposes real costs — lack of **familiarity**, **rough edges**, continual language **changes**. These are **real obstacles**.

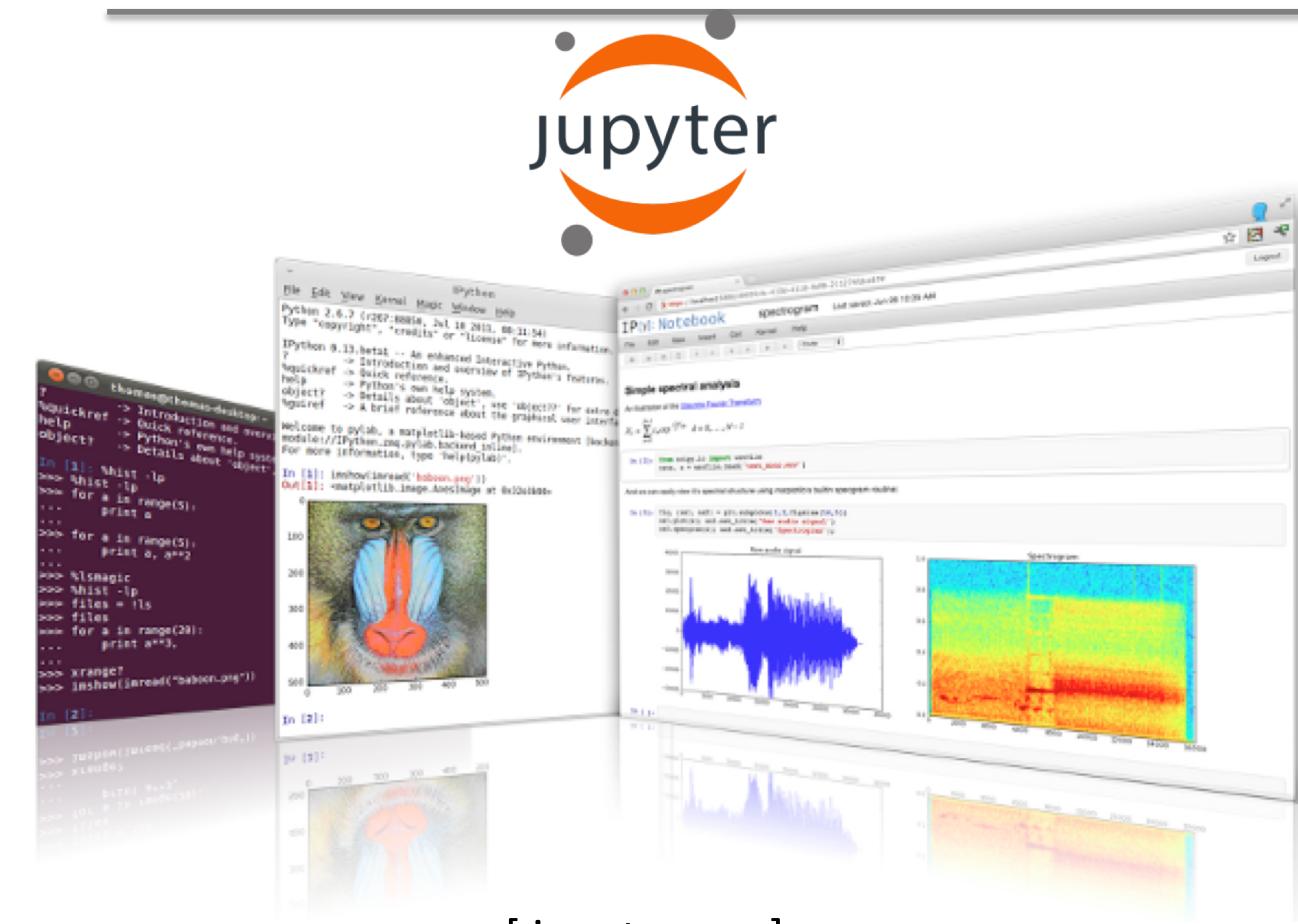
But it also gives you advantages that Matlab/Python users don't have.

But I lose access to all the libraries available for other languages?

Very easy to call C/Fortran libraries from Julia, and also to call Python...

Julia leverages Python...

Directly call Python libraries (PyCall package),
e.g. to plot with Matplotlib (PyPlot package), and also...



[jupyter.org]

via IPython/Jupyter:

Modern multimedia
interactive notebooks
mixing code, results,
graphics, rich text,
equations, interaction

“IJulia”

goto live IJulia notebook demo...

Go to juliabox.org for install-free IJulia on the Amazon cloud

See also julialang.org for more tutorial materials...