

Introduction to Julia: Why are we doing this to you?

(Spring 2020)

Steven G. Johnson, MIT Applied Math

MIT classes 18.06, 18.S191, 18.S096,
18.303, 18.330, 18.08[56],
18.335, 18.337, ...

these slides, cheatsheets, links:

<https://github.com/mitmath/julia-mit>

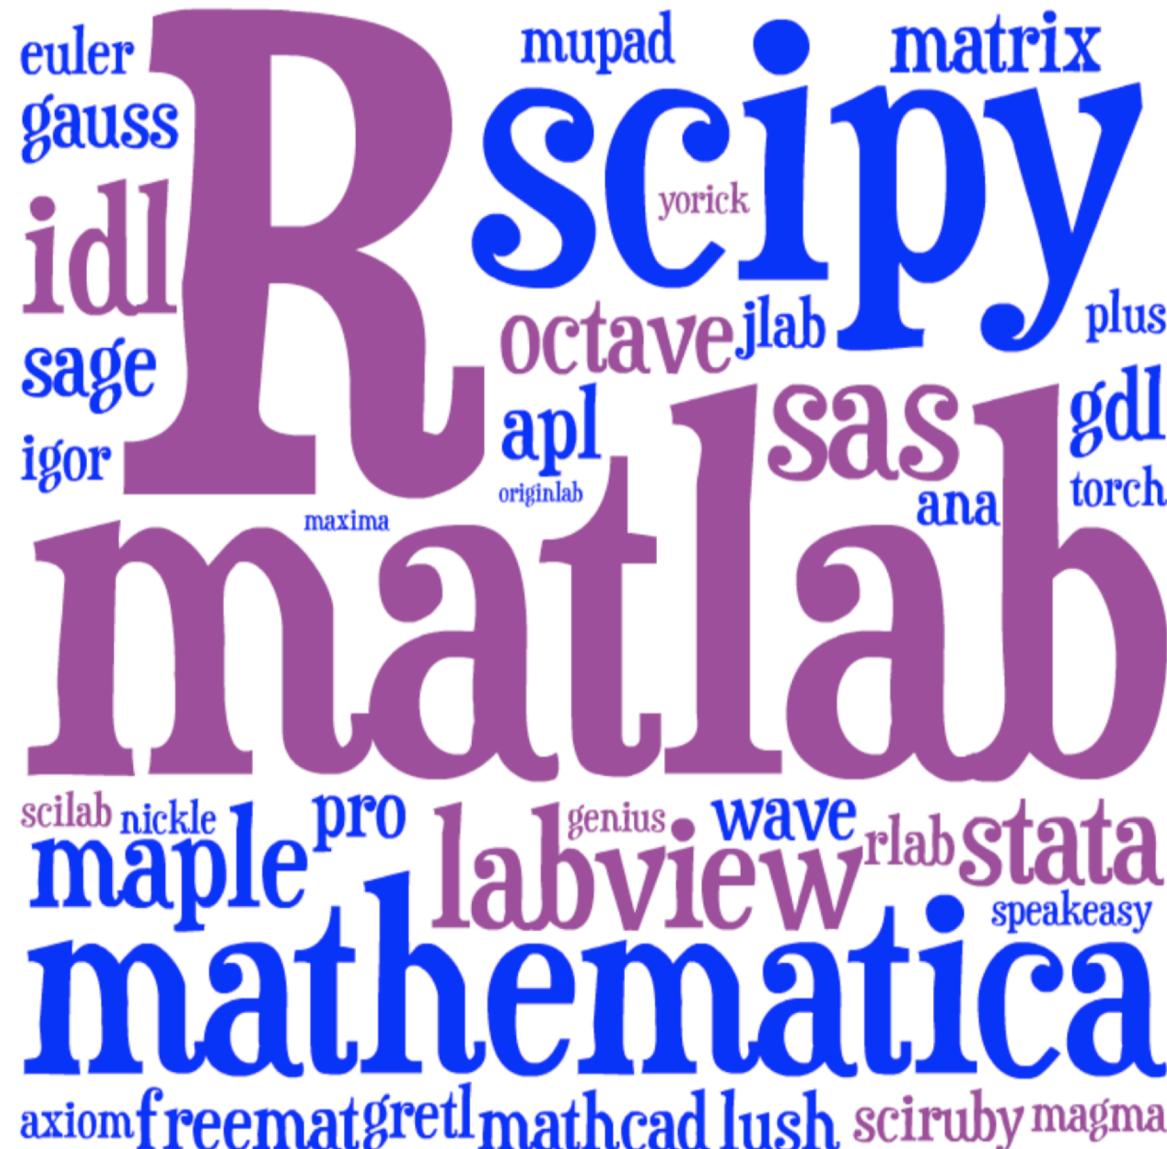
What language for teaching scientific computing?

For the most part, these are not hard-core programming courses, and we only need little “throw-away” scripts and toy numerical experiments.

Almost any high-level, interactive (dynamic) language with easy facilities for linear algebra ($Ax=b$, $Ax=\lambda x$), plotting, mathematical functions, and working with large arrays of data would be fine.

And there are lots of choices...

Lots of choices for interactive math...



[image: Viral Shah]

Just pick the most popular?
Matlab or Python or R?

*We must often use another a language
for our “real work”.*

Traditional HL computing languages
hit a performance wall in “real” work
... eventually force you to C, Cython,
Fortran, Numba...

A new programming language?

Viral Shah

Jeff Bezanson



Stefan Karpinski

Alan Edelman



[MIT]



julialang.org

[begun 2009, “0.1” in 2013, ~50k commits,
1.0 release in 2018, 1.5 in Aug. 2020]

[40+ developers with 100+ commits,
4000+ external packages, 7th JuliaCon in 2020]

As **high-level and interactive** as Matlab or Python+IPython,
as **general-purpose** as Python,
as productive for **technical** work as Matlab or Python+SciPy,
but as **fast as C**.

Lots of performance discussions,
benchmarks, tutorials online...

Let's just look at
one simple example.

Generating Vandermonde matrices

given $x = [\alpha_1, \alpha_2, \dots]$, generate:

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

NumPy (`numpy.vander`): [follow links]

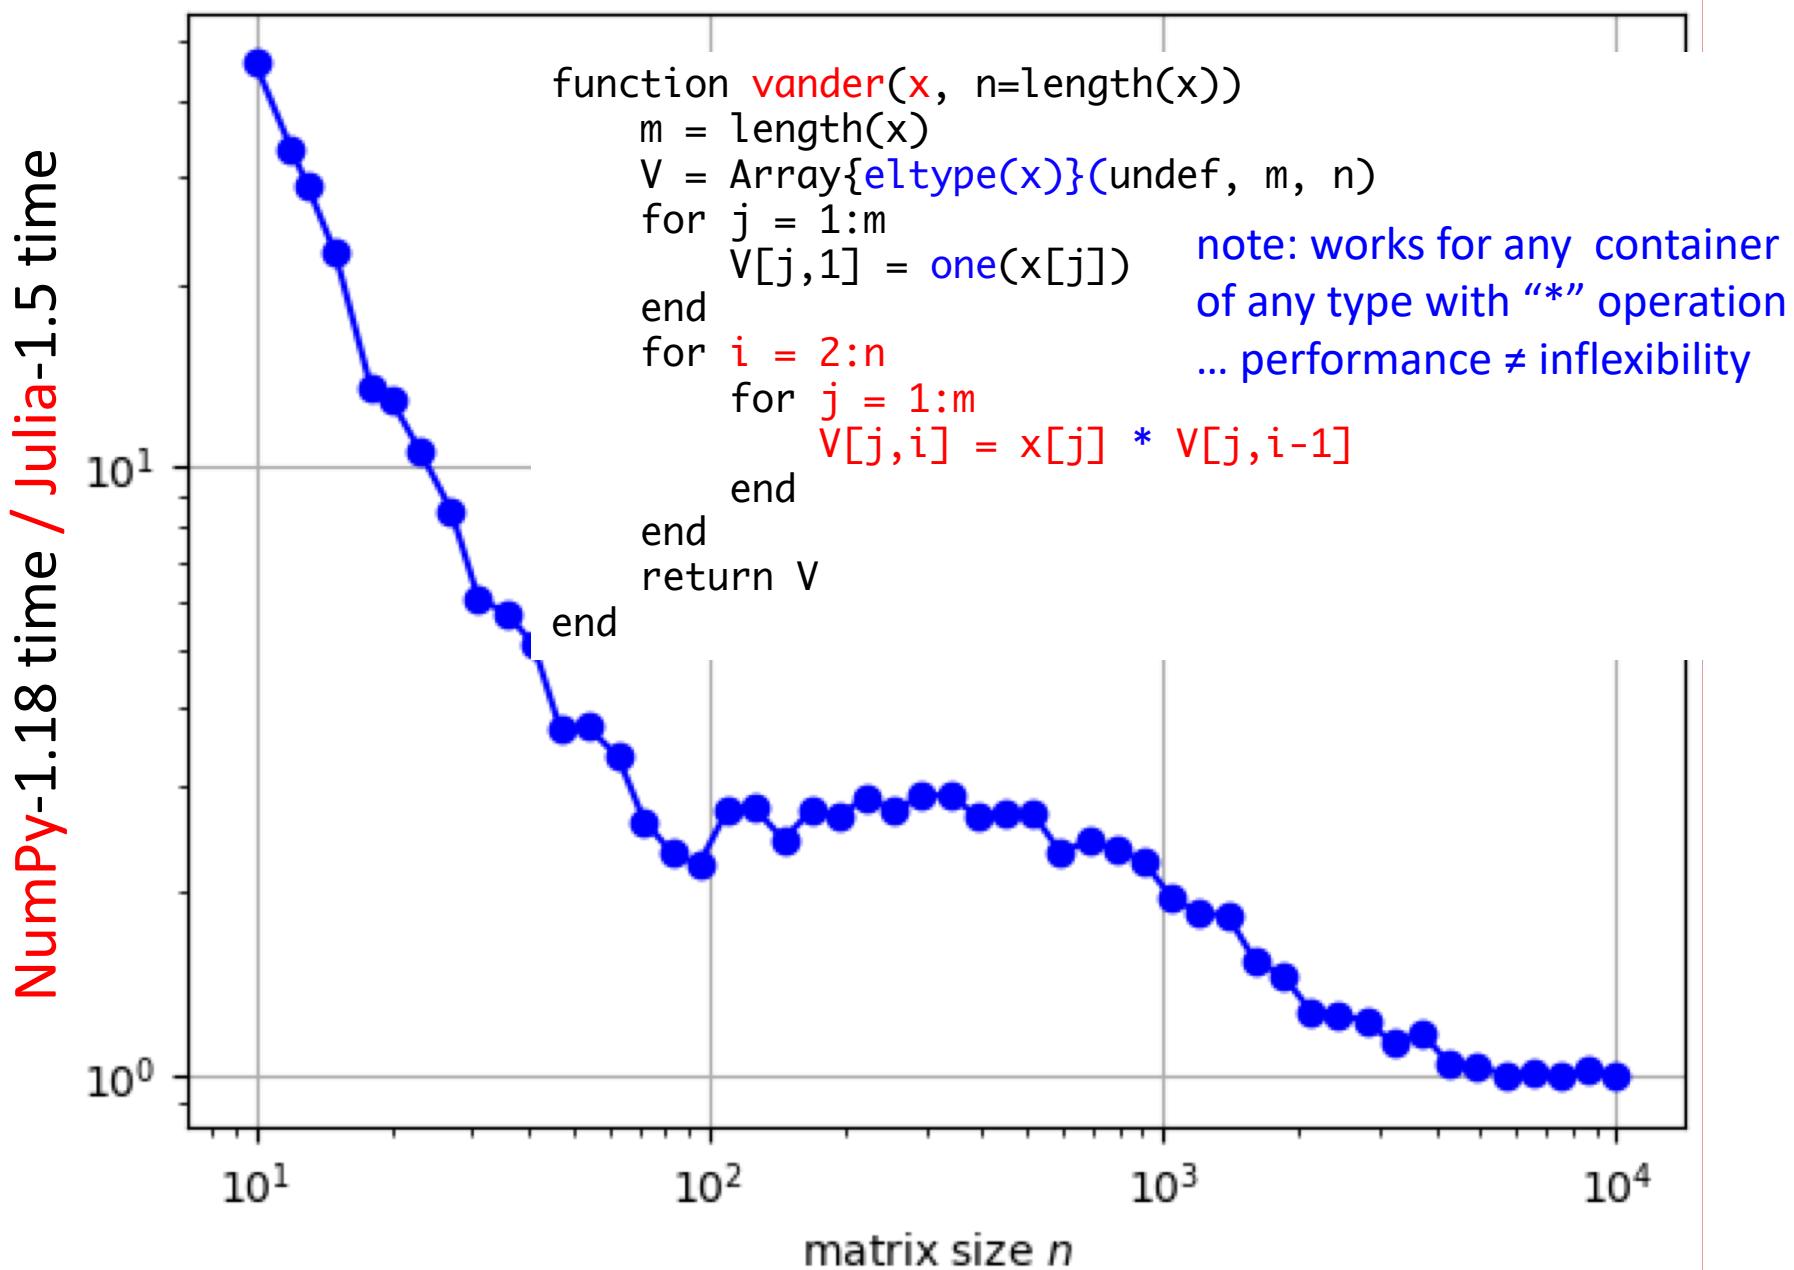
[Python code](#) ...wraps [C code](#)
... wraps [generated C code](#)

type-generic at high-level, but
low level limited to small set of types.

Writing fast code “in” Python or Matlab = mining the standard library
for pre-written functions (implemented in C or Fortran).

If the problem doesn’t “vectorize” into built-in functions,
if you have to write your own inner loops ... sucks for you.

Vandermonde matrices: Julia vs NumPy



Why is Julia fast? A topic for another day.

See my 6.172 “guest lecture” from 2018:

<https://bit.ly/2QUrgB4>

The screenshot shows the MIT OpenCourseWare homepage. At the top, there is a navigation bar with links for "Home", "FIND COURSES", "For Educators", "Give Now", and "About". A red button on the right says "Subscribe to the OCW Newsletter". Below the navigation bar, a breadcrumb trail shows the path: Home > Courses > Electrical Engineering and Computer Science > Performance Engineering of Software Systems > Lecture Videos > Lecture 23: High Performance in Dynamic Languages. The main content area features a large video player for "Lecture 23: High Performance in Dynamic Languages". To the left of the video player is a sidebar with links for "COURSE HOME", "SYLLABUS", "CALENDAR", "READINGS", "LECTURE VIDEOS" (which is highlighted in red), "LECTURE SLIDES", and "ASSIGNMENTS". A small "Interactive Transcript" link is located at the bottom right of the video player.

The original paper:

arXiv.org > cs > arXiv:1411.1607

Computer Science > Mathematical Sc

[Submitted on 6 Nov 2014 (v1), last revised 19 Jul 2015 (this version, v4)]

Julia: A Fresh Approach to Numerical Computing

Jeff Bezanson, Alan Edelman, Stefan Karpinski, Viral B. Shah

But I don't “need” performance!

For lots of problems, especially “toy” problems in courses, Matlab/Python performance is **good enough**.

But if use those languages for all of your “easy” problems, then **you won't be prepared to switch when you hit a hard problem.** When you **need** performance, it is too late.

You **don't** want to learn a **new language** at the **same time** that you are solving **your first truly difficult** computational problem.

Just vectorize your code?

= rely on mature **external libraries**,
operating on **large blocks** of data,
for performance-critical code

Good advice! But...

- Someone has to write those libraries.
- Eventually that person will be **you**.
 - **some problems** are impossible or just very awkward to vectorize.

But everyone else is using
Matlab/Python/R/...

Julia is still a young, niche language. That imposes real costs — lack of **familiarity**, **rough** edges, continual language **changes**. **These are real obstacles.**

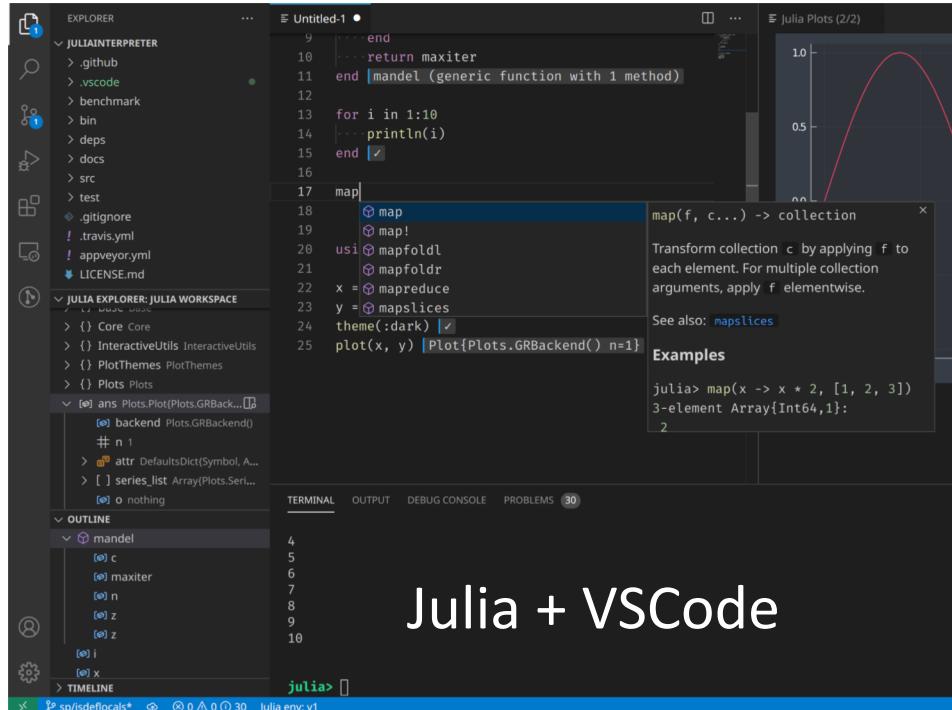
But it also gives you advantages that Matlab/Python users don't have.

But I lose access to all the libraries available for other languages?

Very easy to call C/Fortran libraries from Julia, and also to call Python...

Julia Coding Environments

Of course, full-featured “IDE” coding environments are available with **VSCode**, Atom, not to mention Emacs, Vim, ...



Julia + VSCode

and Julia has the usual features for **organizing large programs** into files, data structures, modules...

+ integrated **hypertext documentation** support...

+ built-in github-integrated **package manager** for distributing & installing packages and dependencies...

... but today we're only looking at **interactive snippets of code**:
Julia as a “glorified calculator”

Julia Interactive-Computing Environments



Jupyter notebooks:
code+text+
equations+images+...

+ Julia backend
= IJulia



Pluto.jl



Fourier sine series

It is a remarkable fact that the sine functions $\sin(n\pi x)$ are **orthogonal** under the "dot" product:

$$\int_0^1 \sin(m\pi x) \sin(n\pi x) dx = \begin{cases} 0 & m \neq n \\ \frac{1}{2} & m = n \end{cases}$$

Let's plot a few of these functions:

In [2]:

```
x = linspace(0,1,200)
for n = 1:4; plot(x, sin.(n*pi*x), "-"); end
legend([L"\sin(\pi x)", L"\sin(2\pi x)", L"\sin(3\pi x)",
title("orthogonal sine functions")]
```

